Field note 001 · Media verification

# A synthetic mountain landscape

As generative models have improved, the question of whether a person can recognise synthetic media by sight, style or sound has moved from curiosity to everyday necessity. Controlled experiments on AI-generated faces, written text and cloned speech have repeatedly found that untrained observers perform only slightly better than chance, and that confidence in a judgement is a poor guide to its accuracy. Automated detectors fare little better in the wild: they tend to degrade when content is compressed, cropped or produced by a model they were not trained on. Researchers and standards bodies have therefore shifted attention from judging appearance to establishing provenance, meaning signed metadata such as content credentials, links to an original publication, corroboration by independent outlets, and a traceable author and date. To illustrate the idea, the verification round in this game presents each player with a fixed set of six items (two images, two short texts, one news report and one audio clip) and asks whether each is authentic or synthetic. Every correct label earns exactly one move, so the number of moves a player carries into the final challenge equals the number of items classified correctly, and the size of the set caps that number. The practical lesson is to check where an item came from before deciding how it looks.

*Illustrative note written for the game. It summarises general findings in the field and does not cite a single study.*